

```
# Data Loading
from google.colab import drive
import pandas as pd
drive.mount('/content/drive')
file_path = '/content/drive/My Drive/diabetes.csv'
df = pd.read_csv(file_path)
```

Mounted at /content/drive

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

Next steps:

[Generate code with df](#)[View recommended plots](#)

```
df.tail()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Pregnancies           768 non-null    int64
 1   Glucose               768 non-null    int64
 2   BloodPressure         768 non-null    int64
 3   SkinThickness         768 non-null    int64
 4   Insulin               768 non-null    int64
 5   BMI                   768 non-null    float64
 6   DiabetesPedigreeFunction 768 non-null    float64
 7   Age                   768 non-null    int64
 8   Outcome               768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
df.columns
```

```

Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')

```

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.00
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.99
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.88
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.00
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.30
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.00
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.60
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.10

```
df.shape
```

```
(768, 9)
```

```

# distribution of outcome variable
df.Outcome.value_counts()*100/len(df)
df['Outcome'].value_counts()*100/len(df)

```

```

Outcome
0    65.104167
1    34.895833


```

Name: count, dtype: float64

```
# plot the hist of the age variable
import matplotlib.pyplot as plt
plt.figure(figsize=(8,7))
plt.xlabel('Age', fontsize=10)
plt.ylabel('Count', fontsize=10)
df['Age'].hist(edgecolor="black")
```

<Axes: xlabel='Age', ylabel='Count'>








```
print("MAX AGE: "+str(df['Age'].max()))
print("MIN AGE: "+str(df['Age'].min()))
```

```
MAX AGE: 81
MIN AGE: 21
```

```
import seaborn as sns
import warnings
warnings.simplefilter(action='ignore')
fig,ax = plt.subplots(4,2, figsize=(20,20))
sns.distplot(df.Pregnancies, bins=20, ax=ax[0,0], color="red")
sns.distplot(df.Glucose, bins=20, ax=ax[0,1], color="red")
sns.distplot(df.BloodPressure, bins=20, ax=ax[1,0], color="red")
sns.distplot(df.SkinThickness, bins=20, ax=ax[1,1], color="red")
sns.distplot(df.Insulin, bins=20, ax=ax[2,0], color="red")
sns.distplot(df.BMI, bins=20, ax=ax[2,1], color="red")
sns.distplot(df.DiabetesPedigreeFunction, bins=20, ax=ax[3,0], color="red")
sns.distplot(df.Age, bins=20, ax=ax[3,1], color="red")
```

```
<Axes: xlabel='Age', ylabel='Density'>
```

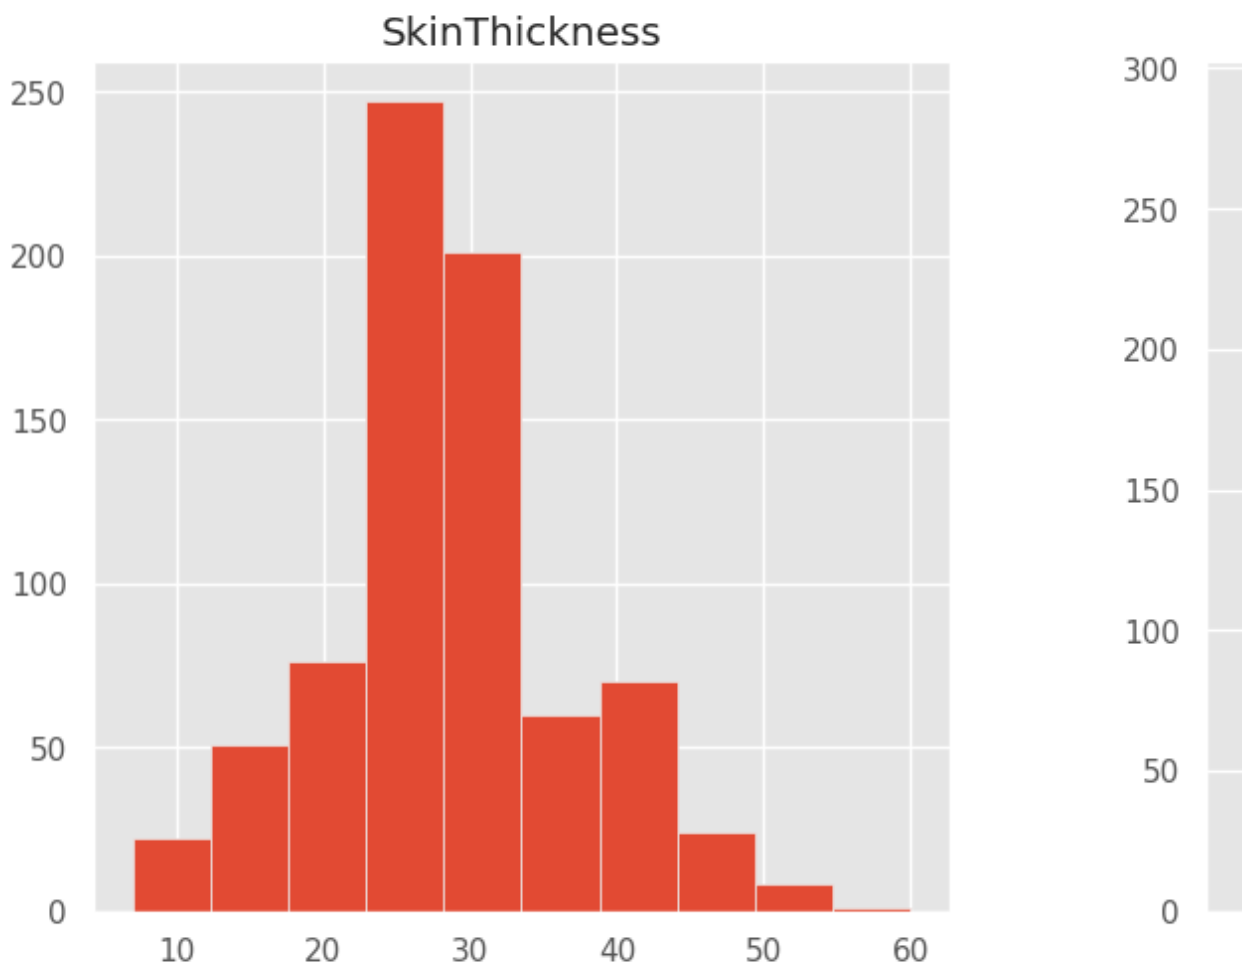






```
df.hist(figsize = (20,20))
```

```
array([[<Axes: title={'center': 'Pregnancies'}>,
<Axes: title={'center': 'Glucose'}>,
<Axes: title={'center': 'BloodPressure'}>],
[<Axes: title={'center': 'SkinThickness'}>,
<Axes: title={'center': 'Insulin'}>,
<Axes: title={'center': 'BMI'}>],
[<Axes: title={'center': 'DiabetesPedigreeFunction'}>,
<Axes: title={'center': 'Age'}>,
<Axes: title={'center': 'Outcome'}>]], dtype=object)
```






DiabetesPedigreeFunction



```
df['Outcome'].value_counts().plot(kind='barh')
```

```
<Axes: ylabel='Outcome'>
```





```
plt.figure(figsize=(5, 5))
sns.countplot(x='Outcome', hue='Outcome', data=df, palette='deep')
plt.title("Outcome")
plt.xlabel("Has Diabetes")
plt.xticks([0, 1], ('False', 'True'))
plt.ylabel("Count")
plt.show()
```



## HAS DIABETES

```
df.groupby("Outcome").agg({'Pregnancies': 'mean'})
```

Pregnancies	
Outcome	
0	3.298000
1	4.865672

```
df.groupby("Outcome").agg({'Pregnancies': 'max'})
```

Pregnancies	
Outcome	
0	13
1	17

```
df.groupby("Outcome").agg({'Glucose': 'mean'})
```

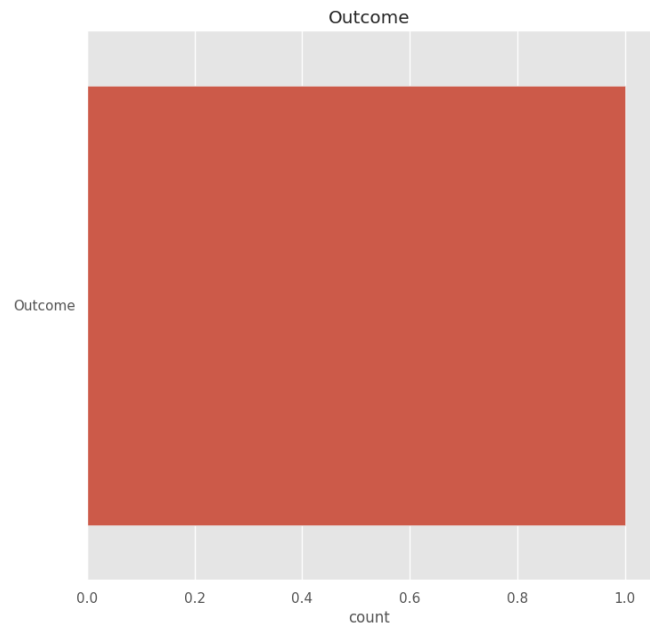

Glucose	
Outcome	
0	109.980000
1	141.257463

```
df.groupby("Outcome").agg({'Glucose': 'max'})
```

Glucose	
Outcome	
0	197
1	199

```
import numpy as np
import statsmodels.api as sm
sns.set()
plt.style.use("ggplot")
%matplotlib inline
```


```
f,ax = plt.subplots(1,2, figsize=(18,8))
df['Outcome'].value_counts().plot.pie(explode=[0,0.1],autopct = "%1.1f%%", ax=f)
ax[0].set_title('target')
ax[0].set_ylabel('')
sns.countplot('Outcome', ax=ax[1])
ax[1].set_title('Outcome')
plt.show()
```




```
df.corr()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	
Pregnancies	1.000000	0.129459	0.141282	-0.081672	(
Glucose	0.129459	1.000000	0.152590	0.057328	(
BloodPressure	0.141282	0.152590	1.000000	0.207371	(
SkinThickness	-0.081672	0.057328	0.207371	1.000000	(
Insulin	-0.073535	0.331357	0.088933	0.436783	.
BMI	0.017683	0.221071	0.281805	0.392573	(
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	(
Age	0.544341	0.263514	0.239528	-0.113970	-(
Outcome	0.221898	0.466581	0.065068	0.074752	(

```
f,ax = plt.subplots(figsize=[20,15])
sns.heatmap(df.corr(), annot=True, fmt = '.2f', ax=ax, cmap='magma')
ax.set_title("Correlation Matrix", fontsize=20)
plt.show()
```





```
# Computer correlation matrix  
corr = df.corr()
```

```
# Generate a mask for the upper triangle
mask = np.zeros_like(corr)
mask[np.triu_indices_from(mask)] = True

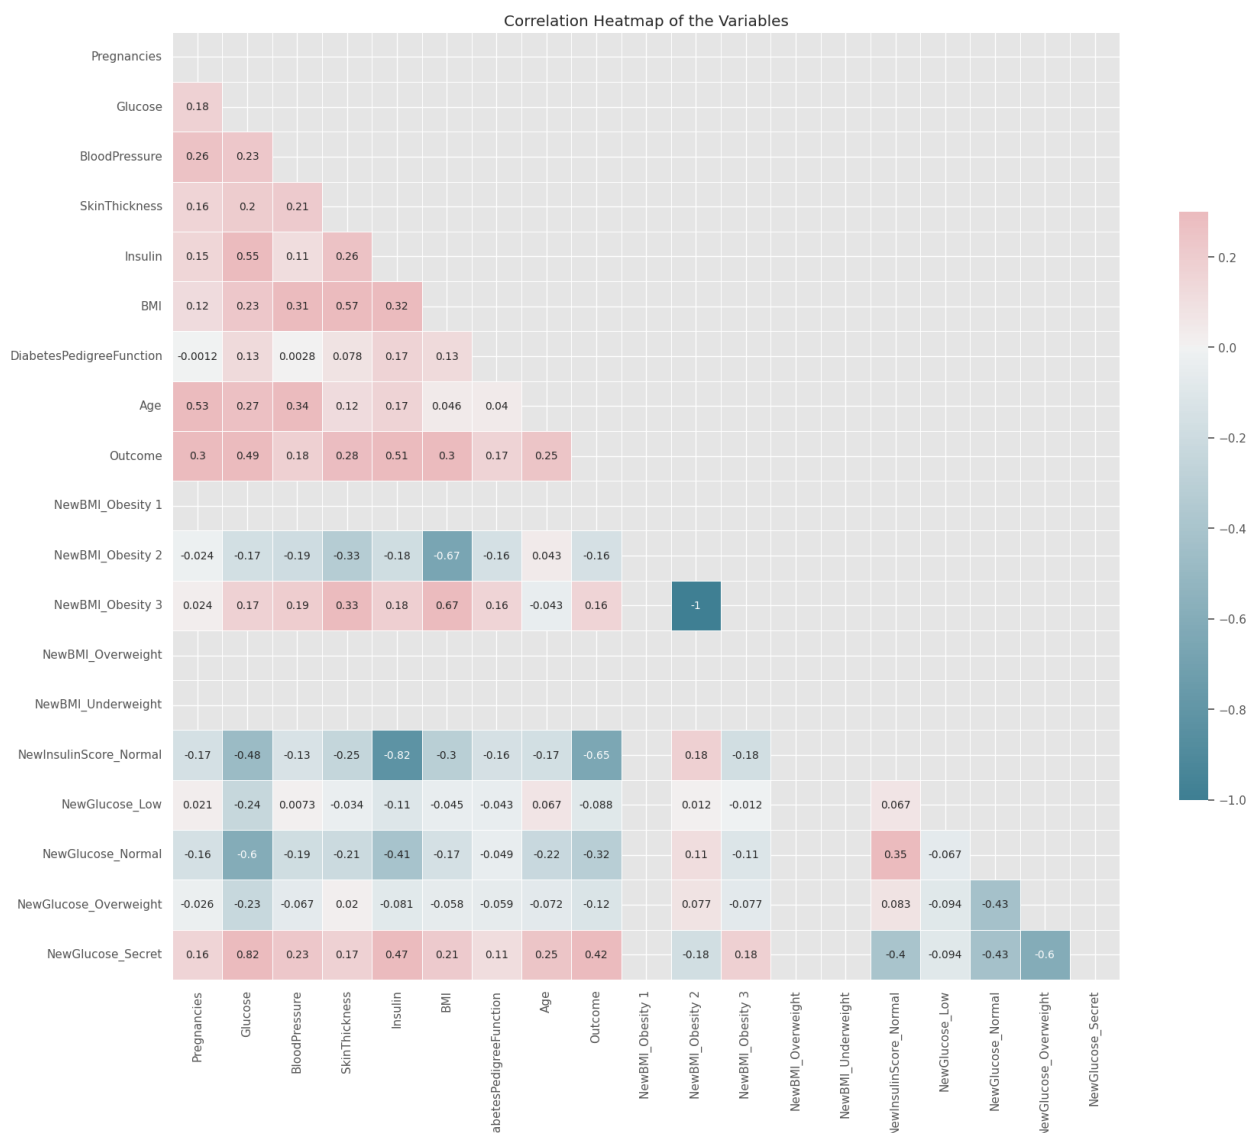
# Set up the matplotlib figure
fig, ax = plt.subplots(figsize=(20, 20))

# Generate a custom diverging colourmap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(
    corr, mask=mask, cmap=cmap, vmax=.3, center=0,
    square=True, linewidths=.5, cbar_kws={"shrink": .5},
    annot=True
)

ax.set_title('Correlation Heatmap of the Variables')

plt.show()
```



```
df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin',  
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
      dtype='object')
```

```
df[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
    'BMI', 'DiabetesPedigreeFunction', 'Age']] = df[['Pregnancies', 'Glucose',  
    'BMI', 'DiabetesPedigreeFunction', 'Age']].replace(0, np.NaN)
```

```
# Data preprocessing Part
```

```
df.isnull().sum()
```

```
Pregnancies      111
Glucose           5
BloodPressure     35
SkinThickness    227
Insulin          374
BMI              11
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```


```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	NaN	33.6	
1	1.0	85.0	66.0	29.0	NaN	26.6	
2	8.0	183.0	64.0	NaN	NaN	23.3	
3	1.0	89.0	66.0	23.0	94.0	28.1	
4	NaN	137.0	40.0	35.0	168.0	43.1	


Next steps:

[Generate code with df](#)[View recommended plots](#)

```
import missingno as msno
msno.bar(df, color="orange")
```







```
# Handling Missing values :
def median_target(var):
    temp = df[df[var].notnull()]
    temp = temp[[var, 'Outcome']].groupby(['Outcome'])[var].median().reset_index()
    return temp
```

```
columns = df.columns
columns = columns.drop("Outcome")
for i in columns:
    median_target(i)
    df.loc[(df['Outcome'] == 0) & (df[i].isnull()), i] = median_target(i)[i][0]
    df.loc[(df['Outcome'] == 1) & (df[i].isnull()), i] = median_target(i)[i][1]
```

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6.0	148.0	72.0	35.0	169.5	33.6	0.3461	33.0	0

1	1.0	85.0	66.0	29.0	102.5	26.6
2	8.0	183.0	64.0	32.0	169.5	23.3
3	1.0	89.0	66.0	23.0	94.0	28.1
4	5.0	137.0	40.0	35.0	168.0	43.1

Next steps:

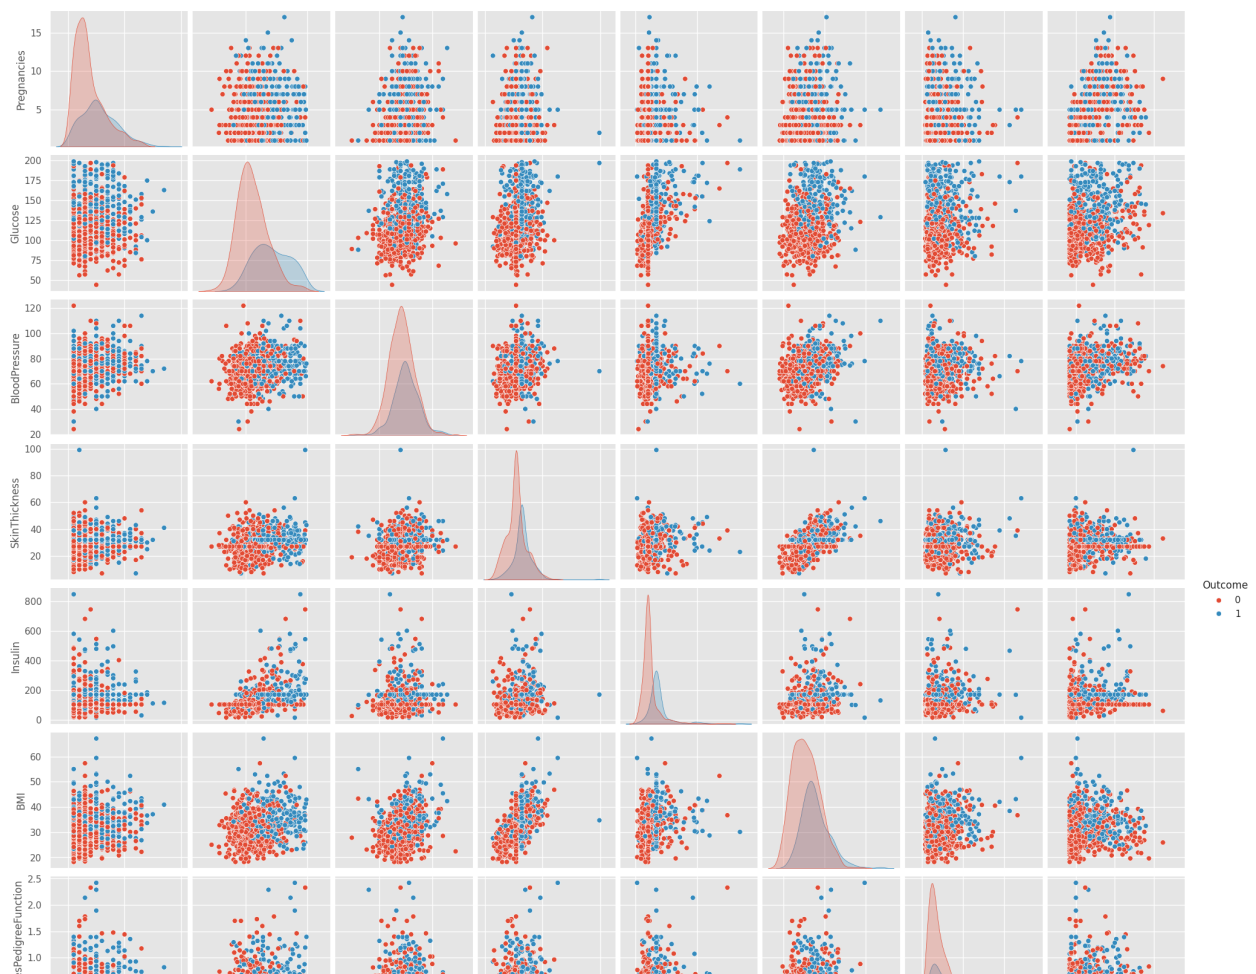
[Generate code with df](#)[View recommended plots](#)


```
df.isnull().sum()
```

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI             0
DiabetesPedigreeFunction  0
Age             0
Outcome          0
dtype: int64
```

# pair plot

```
p = sns.pairplot(df, hue="Outcome")
```





```
import statsmodels.api as sm
# Outlier Detection
# IQR+Q1
# 50%
# 24.65->25%+50%
# 24.65->25%
for feature in df:
    Q1 = df[feature].quantile(0.25)
    Q3 = df[feature].quantile(0.75)
    IQR = Q3-Q1
    lower = Q1 - 1.5*IQR
```

```

lower = Q1-1.5*IQR
upper = Q3+1.5*IQR
if df[(df[feature]>upper)].any(axis=None):
    print(feature, "yes")
else:
    print(feature, "no")

```

```

Pregnancies yes
Glucose no
BloodPressure yes
SkinThickness yes
Insulin yes
BMI yes
DiabetesPedigreeFunction yes
Age yes
Outcome no


```


```

plt.figure(figsize=(8,7))
sns.boxplot(x= df["Insulin"], color="red")

```


<Axes: xlabel='Insulin'>

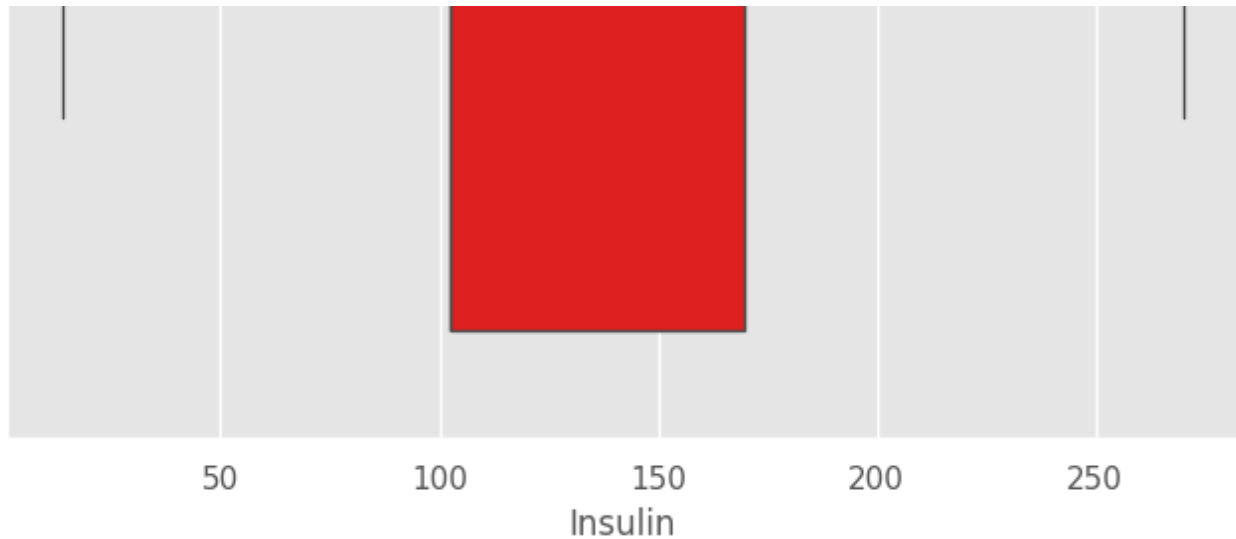




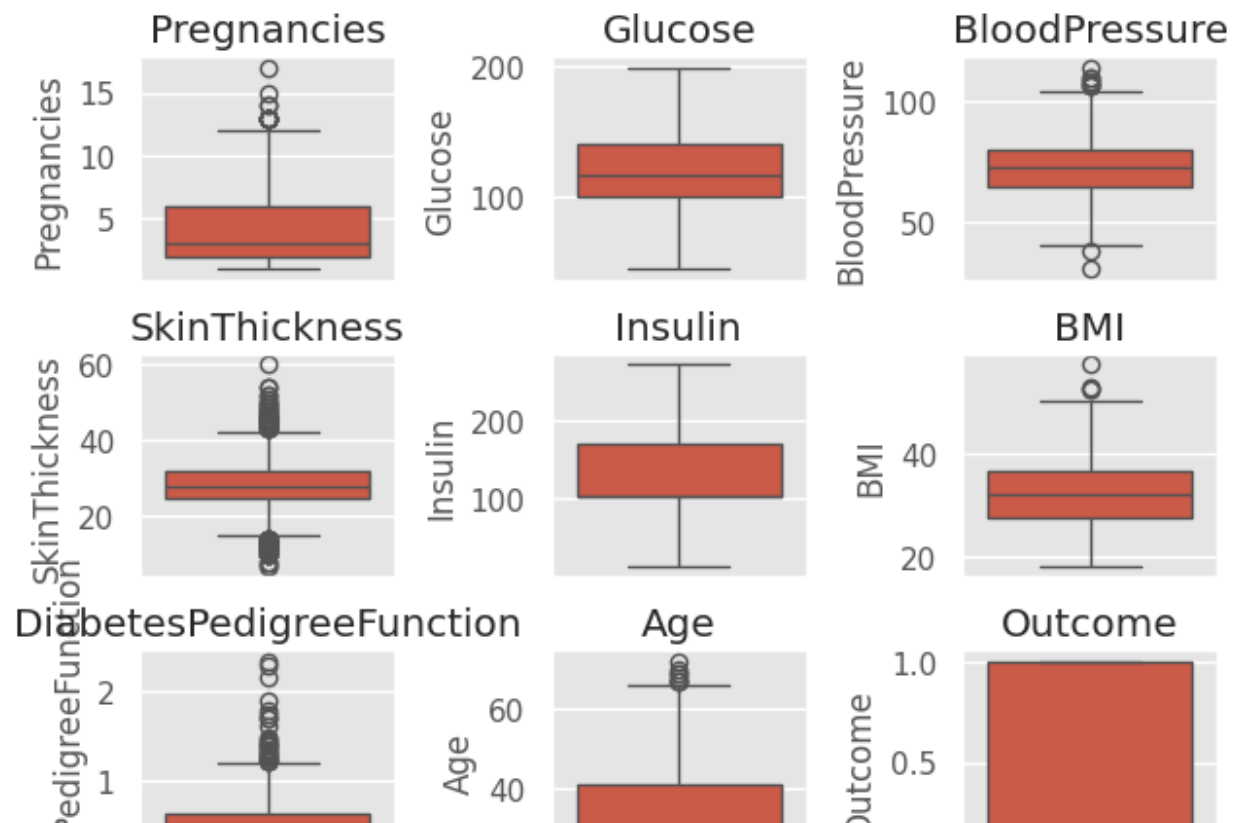
```
Q1 = df.Insulin.quantile(0.25)
Q3 = df.Insulin.quantile(0.75)
IQR = Q3-Q1
lower = Q1-1.5*IQR
upper = Q3+1.5*IQR
df.loc[df['Insulin']>upper, "Insulin"] = upper
plt.figure(figsize=(8,7))
sns.boxplot(x= df["Insulin"], color="red")
```


<Axes: xlabel='Insulin'>





```
repeat = 1
for col in df.select_dtypes(exclude = 'O').columns[:9]:
    plt.subplot(3,3, repeat)
    sns.boxplot(df[col])
    plt.title(col)
    repeat +=1
plt.tight_layout()
plt.show()
```





```
# local outlier factor
from sklearn.neighbors import LocalOutlierFactor
lof = LocalOutlierFactor(n_neighbors=10)
lof.fit_predict(df)
```

[illegible]

```
1, 1, 1, 1, 1, 1, -1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, -1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, -1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, -1, 1, 1, 1, -1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, -1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1])
```

df.head()

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	
1	1.0	85.0	66.0	29.0	102.5	26.6	
2	8.0	183.0	64.0	32.0	169.5	23.3	
3	1.0	89.0	66.0	23.0	94.0	28.1	
4	5.0	137.0	40.0	35.0	168.0	43.1	

Next steps:

[Generate code with df](#)




[View recommended plots](#)



```
plt.figure(figsize=(8,7))
sns.boxplot(x= df["Pregnancies"], color="red")
```

```
<Axes: xlabel='Pregnancies'>
```



```
df_scores = lof.negative_outlier_factor_  
np.sort(df_scores)[0:20]
```

```
array([-3.06509976, -2.38250393, -2.15557018, -2.11501347, -2.08356175,  
       1.05206655, 1.02550204, 1.74074227, 1.7220214, 1.71017160,       1.69817160, 1.68617160, 1.67417160, 1.66217160, 1.65017160,  
       1.63817160, 1.62617160, 1.61417160, 1.60217160, 1.59017160])
```

```
-1.95580055, -1.85559384, -1.74974237, -1.7350214 , -1.71017108,  
-1.70215105, -1.68722889, -1.64294601, -1.64180205, -1.61181746,  
-1.61067772, -1.60925053, -1.60214364, -1.59998552, -1.58761193])
```

```
threshold = np.sort(df_scores)[7]  
threshold
```

```
-1.7497423670960557
```

```
outlier = df_scores > threshold  
df = df[outlier]
```

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	1
1	1.0	85.0	66.0	29.0	102.5	26.6	0
2	8.0	183.0	64.0	32.0	169.5	23.3	1
3	1.0	89.0	66.0	23.0	94.0	28.1	0
4	5.0	137.0	40.0	35.0	168.0	43.1	1

Next steps:


[Generate code with df](#)[View recommended plots](#)

```
df.shape
```

```
(760, 9)
```

```
plt.figure(figsize=(8,7))  
sns.boxplot(x= df["Pregnancies"], color="red")
```

```
<Axes: xlabel='Pregnancies'>
```



```
# Feature Engineering
```

```
# Feature Engineering
NewBMI = pd.Series(["Underweight","Normal", "Overweight","Obesity 1", "Obesity 2", "Obesity 3"])
NewBMI
0    Underweight
1         Normal
2    Overweight
3    Obesity 1
4    Obesity 2
5    Obesity 3
dtype: category
Categories (6, object): ['Normal', 'Obesity 1', 'Obesity 2', 'Obesity 3', 'Overweight', 'Underweight']
```

```
df['NewBMI'] = NewBMI
df.loc[df["BMI"]<18.5, "NewBMI"] = NewBMI[0]
df.loc[(df["BMI"]>18.5) & df["BMI"]<=24.9, "NewBMI"] = NewBMI[1]
df.loc[(df["BMI"]>24.9) & df["BMI"]<=29.9, "NewBMI"] = NewBMI[2]
df.loc[(df["BMI"]>29.9) & df["BMI"]<=34.9, "NewBMI"] = NewBMI[3]
df.loc[(df["BMI"]>34.9) & df["BMI"]<=39.9, "NewBMI"] = NewBMI[4]
df.loc[df["BMI"]>39.9, "NewBMI"] = NewBMI[5]
```

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	1
1	1.0	85.0	66.0	29.0	102.5	26.6	0
2	8.0	183.0	64.0	32.0	169.5	23.3	1

Next steps:

[Generate code with df](#)
[View recommended plots](#)

```
# if insulin>=16 & insuline<=166->normal
def set_insuline(row):
    if row["Insulin"]>=16 and row["Insulin"]<=166:
        return "Normal"
    else:
        return "Abnormal"

df = df.assign(NewInsulinScore=df.apply(set_insuline, axis=1))
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
--	-------------	---------	---------------	---------------	---------	-----	----------

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	
1	1.0	85.0	66.0	29.0	102.5	26.6	
2	8.0	183.0	64.0	32.0	169.5	23.3	

Next steps:

[Generate code with df](#)[View recommended plots](#)

```
# Some intervals were determined according to the glucose variable and these were used to create the NewGlucose variable
NewGlucose = pd.Series(["Low", "Normal", "Overweight", "Secret", "High"], dtype=object)
df["NewGlucose"] = NewGlucose
df.loc[df["Glucose"] <= 70, "NewGlucose"] = NewGlucose[0]
df.loc[(df["Glucose"] > 70) & (df["Glucose"] <= 99), "NewGlucose"] = NewGlucose[1]
df.loc[(df["Glucose"] > 99) & (df["Glucose"] <= 126), "NewGlucose"] = NewGlucose[2]
df.loc[df["Glucose"] > 126, "NewGlucose"] = NewGlucose[3]
```

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	
1	1.0	85.0	66.0	29.0	102.5	26.6	
2	8.0	183.0	64.0	32.0	169.5	23.3	

Next steps:

[Generate code with df](#)[View recommended plots](#)

```
# One hot encoding
df = pd.get_dummies(df, columns = ["NewBMI", "NewInsulinScore", "NewGlucose"],
```

```
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	
1	1.0	85.0	66.0	29.0	102.5	26.6	

2	8.0	183.0	64.0	32.0	169.5	23.3
3	1.0	89.0	66.0	23.0	94.0	28.1
4	5.0	137.0	40.0	35.0	168.0	43.1

Next steps:

[Generate code with df](#)[View recommended plots](#)

df.columns

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
      'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome',
      'NewBMI_Obesity 1',
      'NewBMI_Obesity 2', 'NewBMI_Obesity 3', 'NewBMI_Overweight',
      'NewBMI_Underweight', 'NewInsulinScore_Normal', 'NewGlucose_Low',
      'NewGlucose_Normal', 'NewGlucose_Overweight', 'NewGlucose_Secret'],
      dtype='object')
```

```
categorical_df = df[['NewBMI_Obesity 1',
                    'NewBMI_Obesity 2', 'NewBMI_Obesity 3', 'NewBMI_Overweight',
                    'NewBMI_Underweight', 'NewInsulinScore_Normal', 'NewGlucose_Low',
                    'NewGlucose_Normal', 'NewGlucose_Overweight', 'NewGlucose_Secret']]
```

categorical\_df.head()

	NewBMI_Obesity 1	NewBMI_Obesity 2	NewBMI_Obesity 3	NewBMI_Overweight	NewBMI
0	False	True	False	False	
1	False	True	False	False	
2	False	True	False	False	
3	False	True	False	False	

4

False

False

True

False

Next steps:

[Generate code with categorical\\_df](#)[View recommended plots](#)

```

y=df['Outcome']
X=df.drop(['Outcome', 'NewBMI_Obesity 1',
          'NewBMI_Obesity 2', 'NewBMI_Obesity 3', 'NewBMI_Overweight',
          'NewBMI_Underweight', 'NewInsulinScore_Normal', 'NewGlucose_Low',
          'NewGlucose_Normal', 'NewGlucose_Overweight', 'NewGlucose_Secret'], axis=

```

```

cols = X.columns
index = X.index

```

```
X.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6.0	148.0	72.0	35.0	169.5	33.6	
1	1.0	85.0	66.0	29.0	102.5	26.6	
2	8.0	183.0	64.0	32.0	169.5	23.3	
3	1.0	89.0	66.0	23.0	94.0	28.1	
4	5.0	137.0	40.0	35.0	168.0	43.1	

Next steps:

[Generate code with X](#)[View recommended plots](#)

```

from sklearn.preprocessing import RobustScaler
transformer = RobustScaler().fit(X)
X=transformer.transform(X)
X=pd.DataFrame(X, columns = cols, index = index)

```

```
X.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	0.75	0.775	0.000	1.000000	1.000000	0.177778	
1	-0.50	-0.800	-0.375	0.142857	0.000000	-0.600000	
2	1.25	1.650	-0.500	0.571429	1.000000	-0.966667	
3	-0.50	-0.700	-0.375	-0.714286	-0.126866	-0.433333	

	0.50	0.750	-0.375	0.714286	0.120000	0.700000
4	0.50	0.500	-2.000	1.000000	0.977612	1.233333

Next steps:

[Generate code with X](#)[View recommended plots](#)

```
X = pd.concat([X, categorical_df], axis=1)
```

```
X.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	0.75	0.775	0.000	1.000000	1.000000	0.177778	
1	-0.50	-0.800	-0.375	0.142857	0.000000	-0.600000	
2	1.25	1.650	-0.500	0.571429	1.000000	-0.966667	
3	-0.50	-0.700	-0.375	-0.714286	-0.126866	-0.433333	
4	0.50	0.500	-2.000	1.000000	0.977612	1.233333	

5 rows x 28 columns

```
from sklearn.preprocessing import scale, StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score, mean_squared_error
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import KFold
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
```



▼ LogisticRegression

LogisticRegression()

```
y_pred = log_reg.predict(X_test)
accuracy_score(y_train, log_reg.predict(X_train))
```

```
0.8470394736842105
```

```
log_reg_acc = accuracy_score(y_test, log_reg.predict(X_test))
confusion_matrix(y_test, y_pred)
```

```
array([[88, 10],
       [ 6, 48]])
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	98
1	0.83	0.89	0.86	54
accuracy			0.89	152
macro avg	0.88	0.89	0.89	152
weighted avg	0.90	0.89	0.90	152

#KNN-classifier approach:

```
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print(accuracy_score(y_train, knn.predict(X_train)))
knn_acc = accuracy_score(y_test, knn.predict(X_test))
print(accuracy_score(y_test, knn.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```


```
0.8782894736842105
```

```
0.8947368421052632
```

```
[[88 10]
 [ 6 48]]
```

	precision	recall	f1-score	support
0	0.94	0.90	0.92	98
1	0.83	0.89	0.86	54
accuracy			0.89	152
macro avg	0.88	0.89	0.89	152
weighted avg	0.90	0.89	0.90	152

```
# SVM-classifier approach:
svc = SVC(probability=True)
parameter = {
    "gamma":[0.0001, 0.001, 0.01, 0.1],
    'C': [0.01, 0.05,0.5, 0.01, 1, 10, 15, 20]
}
grid_search = GridSearchCV(svc, parameter)
grid_search.fit(X_train, y_train)
```



```
# best_parameter
grid_search.best_params_

{'C': 1, 'gamma': 0.1}
```

```
grid_search.best_score_

0.8602086438152012
```

```
svc = SVC(C=10, gamma = 0.01, probability=True)
svc.fit(X_train, y_train)
y_pred = svc.predict(X_test)
print(accuracy_score(y_train, svc.predict(X_train)))
svc_acc = accuracy_score(y_test, svc.predict(X_test))
print(accuracy_score(y_test, svc.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

0.8848684210526315
0.9078947368421053
[[90  8]
 [ 6 48]]
```

	precision	recall	f1-score	support
0	0.94	0.92	0.93	98
1	0.86	0.89	0.87	54
accuracy			0.91	152
macro avg	0.90	0.90	0.90	152
weighted avg	0.91	0.91	0.91	152

```
# Decision Tree Approach:
DT = DecisionTreeClassifier()
DT.fit(X_train, y_train)
y_pred = DT.predict(X_test)
print(accuracy_score(y_train, DT.predict(X_train)))

print(accuracy_score(y_test, DT.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

1.0  
0.868421052631579  
[[85 13]  
 [ 7 47]]

		precision	recall	f1-score	support
	0	0.92	0.87	0.89	98
	1	0.78	0.87	0.82	54
	accuracy			0.87	152
	macro avg	0.85	0.87	0.86	152
	weighted avg	0.87	0.87	0.87	152

```
# hyperparameter tuning of Decision Theoretic Model:
grid_param = {
    'criterion':['gini','entropy'],
    'max_depth' : [3,5,7,10],
    'splitter' : ['best','radom'],
    'min_samples_leaf':[1,2,3,5,7],
    'min_samples_split':[1,2,3,5,7],
    'max_features':['auto','sqrt','log2']
}
grid_search_dt = GridSearchCV(DT, grid_param, cv=50, n_jobs=-1, verbose = 1)
grid_search_dt.fit(X_train, y_train)
```

Fitting 50 folds for each of 1200 candidates, totalling 60000 fits

```
▼ GridSearchCV
GridSearchCV(cv=50, estimator=DecisionTreeClassifier(), n_jobs=-1,
             param_grid={'criterion': ['gini', 'entropy'],
                          'max_depth': [3, 5, 7, 10],
                          'max_features': ['auto', 'sqrt', 'log2'],
                          'min_samples_leaf': [1, 2, 3, 5, 7],
                          'min_samples_split': [1, 2, 3, 5, 7],
                          'splitter': ['best', 'radom']}},
             verbose=1)
  ▼ estimator: DecisionTreeClassifier
    DecisionTreeClassifier()
      ▼ DecisionTreeClassifier
        DecisionTreeClassifier()
```

```
grid_search_dt.best_params_
```

```
{'criterion': 'gini',
 'max_depth': 5,
```

```
'max_features': 'sqrt',
'min_samples_leaf': 7,
'min_samples_split': 2,
'splitter': 'best'}
```

```
grid_search_dt.best_score_
```

```
0.8741025641025643
```

```
DT = grid_search_dt.best_estimator_
y_pred = DT.predict(X_test)
print(accuracy_score(y_train, DT.predict(X_train)))
dt_acc = accuracy_score(y_test, DT.predict(X_test))
print(accuracy_score(y_test, DT.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.8601973684210527
```

```
0.8355263157894737
```

```
[[78 20]
```

```
 [ 5 49]]
```

	precision	recall	f1-score	support
0	0.94	0.80	0.86	98
1	0.71	0.91	0.80	54
accuracy			0.84	152
macro avg	0.82	0.85	0.83	152
weighted avg	0.86	0.84	0.84	152

```
# Random Forest Classifier Approach:
```

```
rand_clf = RandomForestClassifier(criterion = 'entropy', max_depth = 15, max_fe
rand_clf.fit(X_train, y_train)
```

```
▼                                RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=15, max_features=0.75,
                        min_samples_leaf=2, min_samples_split=3,
                        n_estimators=130)
```

```
y_pred = rand_clf.predict(X_test)
```

```

y_pred = rand_clf.predict(X_test)
print(accuracy_score(y_train, rand_clf.predict(X_train)))
rand_acc = accuracy_score(y_test, rand_clf.predict(X_test))
print(accuracy_score(y_test, rand_clf.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

```

```
0.9901315789473685
```

```
0.9013157894736842
```

```
[[89  9]
```

```
 [ 6 48]]
```

	precision	recall	f1-score	support
0	0.94	0.91	0.92	98
1	0.84	0.89	0.86	54
accuracy			0.90	152
macro avg	0.89	0.90	0.89	152
weighted avg	0.90	0.90	0.90	152

```
gbc = GradientBoostingClassifier()
```

```

parameters = {
    'loss': ['deviance', 'exponential'],
    'learning_rate': [0.001, 0.1, 1, 10],
    'n_estimators': [100, 150, 180, 200]
}

```

```

grid_search_gbc = GridSearchCV(gbc, parameters, cv = 10, n_jobs = -1, verbose =
grid_search_gbc.fit(X_train, y_train)

```

Fitting 10 folds for each of 32 candidates, totalling 320 fits

```

▼ GridSearchCV
GridSearchCV(cv=10, estimator=GradientBoostingClassifier(), n_jobs=-1,
param_grid={'learning_rate': [0.001, 0.1, 1, 10],
'loss': ['deviance', 'exponential'],
'n_estimators': [100, 150, 180, 200]},
verbose=1)
  ▼ estimator: GradientBoostingClassifier
  GradientBoostingClassifier()

```

```
▼ GradientBoostingClassifier
GradientBoostingClassifier()
```

```
grid_search_gbc.best_params_
```

```
{'learning_rate': 0.1, 'loss': 'exponential', 'n_estimators': 150}
```

```
grid_search_gbc.best_score_
```

```
0.888032786885246
```

```
gbc = GradientBoostingClassifier(learning_rate = 0.1, loss = 'exponential', n_estimators=150)
gbc.fit(X_train, y_train)
```

```
▼ GradientBoostingClassifier
GradientBoostingClassifier(loss='exponential', n_estimators=150)
```

```
gbc = grid_search_gbc.best_estimator_
y_pred = gbc.predict(X_test)
print(accuracy_score(y_train, gbc.predict(X_train)))
gbc_acc = accuracy_score(y_test, gbc.predict(X_test))
print(accuracy_score(y_test, gbc.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.9983552631578947
0.9144736842105263
[[91  7]
 [ 6 48]]
```

	precision	recall	f1-score	support
0	0.94	0.93	0.93	98
1	0.87	0.89	0.88	54
accuracy			0.91	152
macro avg	0.91	0.91	0.91	152
weighted avg	0.91	0.91	0.91	152

```
from xgboost import XGBClassifier
xgb = XGBClassifier(objective = 'binary:logistic', learning_rate = 0.01, max_de

xgb.fit(X_train, y_train)
```

```
▼ XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytrees=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.01, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=10, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=180, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...)
```

```
y_pred = xgb.predict(X_test)
print(accuracy_score(y_train, xgb.predict(X_train)))
xgb_acc = accuracy_score(y_test, xgb.predict(X_test))
print(accuracy_score(y_test, xgb.predict(X_test)))
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
0.9736842105263158
```

```
0.881578947368421
```



```
[[89  9]
```

```
 [ 9 45]]
```

	precision	recall	f1-score	support
0	0.91	0.91	0.91	98
1	0.83	0.83	0.83	54
accuracy			0.88	152
macro avg	0.87	0.87	0.87	152
weighted avg	0.88	0.88	0.88	152



```
# Model Comparison
models = pd.DataFrame({
    'Model': ['Logistic Regression', 'KNN', 'SVM', 'Decision Tree Classifier',
    'Score': [100*round(log_reg_acc,4), 100*round(knn_acc,4), 100*round(svc_acc,4),
    100*round(gbc_acc,4), 100*round(xgb_acc,4)]
})
models.sort_values(by = 'Score', ascending = False)
```

	Model	Score	
5	Gradient Boosting Classifier	91.45	
2	SVM	90.79	
4	Random Forest Classifier	90.13	
0	Logistic Regression	89.47	
1	KNN	89.47	
6	XgBoost	88.16	
3	Decision Tree Classifier	83.55	

```
import pickle
model = gbc_acc
pickle.dump(model, open("diabetes.pkl", 'wb'))
```


```
from sklearn import metrics
plt.figure(figsize=(16,10))
models = [
{
    'label': 'LR',
    'model': log_reg,
},
{
    'label': 'DT',
    'model': DT,
},
{
    'label': 'SVM',
    'model': svc,
},
{
    'label': 'KNN',
    'model': knn,
},
{
    'label': 'XGBoost',
    'model': xgb,
},
]
```

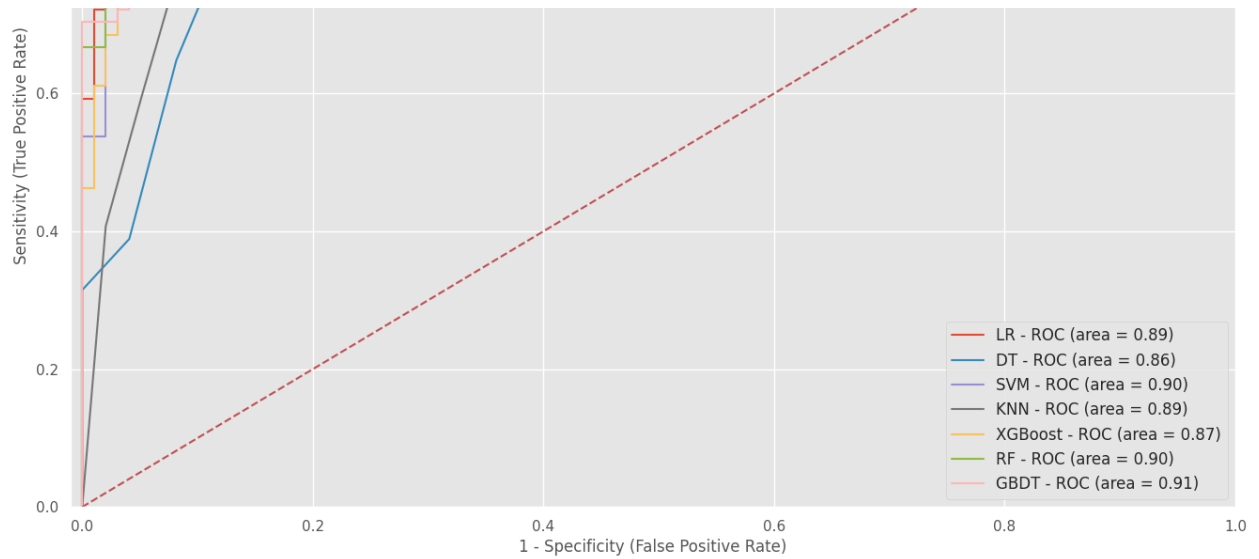
```

    'label': 'RF',
    'model': rand_clf,
},
{
    'label': 'GBDT',
    'model': gbc,
}
]
for m in models:
    model = m['model']
    model.fit(X_train, y_train)
    y_pred=model.predict(X_test)
    fpr1, tpr1, thresholds = metrics.roc_curve(y_test, model.predict_proba(X_te
    auc = metrics.roc_auc_score(y_test,model.predict(X_test))
    plt.plot(fpr1, tpr1, label='%s - ROC (area = %0.2f)' % (m['label'], auc))

plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([-0.01, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('1 - Specificity (False Positive Rate)', fontsize=12)
plt.ylabel('Sensitivity (True Positive Rate)', fontsize=12)
plt.title('ROC - Diabetes Prediction', fontsize=12)
plt.legend(loc="lower right", fontsize=12)
plt.savefig("roc_diabetes.jpeg", format='jpeg', dpi=400, bbox_inches='tight')
plt.show()

```





```
from sklearn import metrics

import numpy as np
import matplotlib.pyplot as plt

models = [
    {
        'label': 'LR',
        'model': log_reg,
    },
    {
        'label': 'DT',
        'model': DT,
    },
]
```

```

- ,
{
    'label': 'SVM',
    'model': svc,
},
{
    'label': 'KNN',
    'model': knn,
},
{
    'label': 'XGBoost',
    'model': xgb,
},
{
    'label': 'RF',
    'model': rand_clf,
},
{
    'label': 'GBDT',
    'model': gbc,
}
]

means_roc = []
means_accuracy = [100*round(log_reg_acc,4), 100*round(dt_acc,4), 100*round(svc_
                    100*round(rand_acc,4), 100*round(gbc_acc,4)]

for m in models:
    model = m['model']
    model.fit(X_train, y_train)
    y_pred=model.predict(X_test)
    fpr1, tpr1, thresholds = metrics.roc_curve(y_test, model.predict_proba(X_te
    auc = metrics.roc_auc_score(y_test,model.predict(X_test))
    auc = 100*round(auc,4)
    means_roc.append(auc)

print(means_accuracy)
print(means_roc)

# data to plot
n_groups = 7
means_accuracy = tuple(means_accuracy)
means_roc = tuple(means_roc)

# create plot
fig, ax = plt.subplots(figsize=(16,10))
index = np.arange(n_groups)
bar_width = 0.35
opacity = 0.8

rects1 = plt.bar(index, means_accuracy, bar_width,
alpha=opacity,
color='mediumslateblue')

```

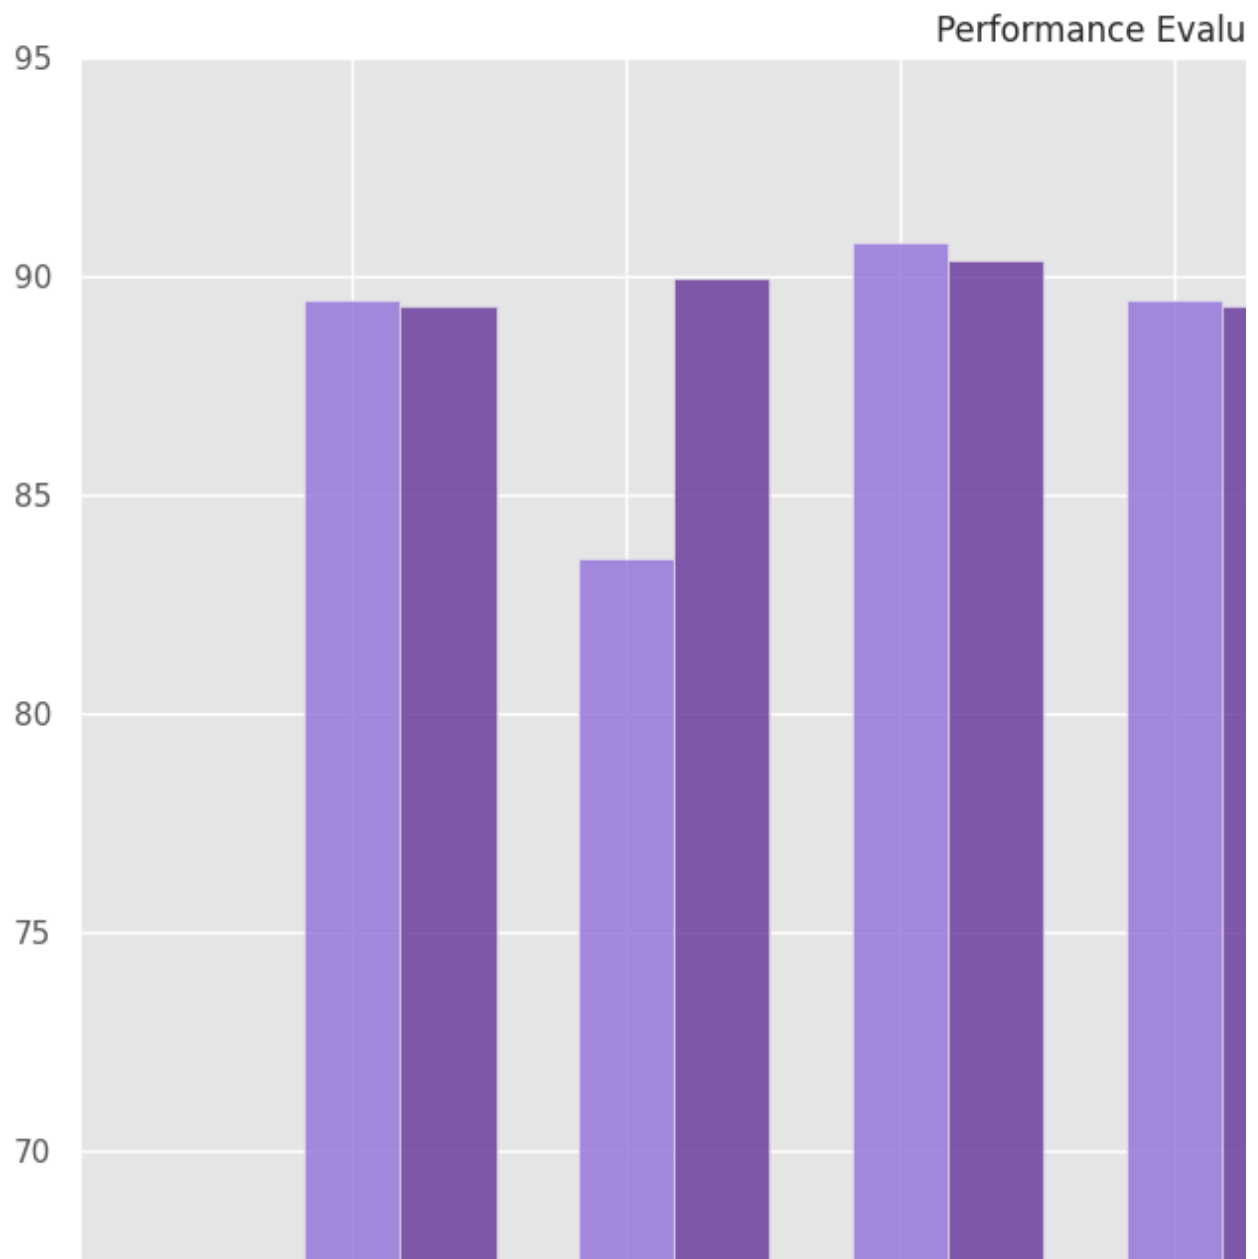
```
color='mediumpurple',
label='Accuracy (%)')
```


```
rects2 = plt.bar(index + bar_width, means_roc, bar_width,
alpha=opacity,
color='rebeccapurple',
label='ROC (%)')
```

```
plt.xlim([-1, 8])
plt.ylim([60, 95])
```

```
plt.title('Performance Evaluation - Diabetes Prediction', fontsize=12)
plt.xticks(index, (' LR', ' DT', ' SVM', ' KNN', 'XGBoost', ' RF', '
plt.legend(loc="upper right", fontsize=10)
plt.savefig("PE_diabetes.jpeg", format='jpeg', dpi=400, bbox_inches='tight')
plt.show()
```

```
[89.47, 83.55, 90.79, 89.47, 88.16000000000001, 90.13, 91.45]
[89.34, 89.95, 90.36, 89.34, 87.07000000000001, 88.92999999999999, 90.86999
```





```
# Statistical Learning plus analysis:
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

def plot_radar(df, bins, column,data):
    # SET DATA
    data_counts = pd.crosstab(pd.cut(df[column], bins=bins), df[data])

    # CREATE BACKGROUND
    datas = set(pd.cut(df[column], bins=bins))

    # Angle of each axis in the plot
    angles = [(n / len(datas)) * 2 * np.pi for n in range(len(datas)+1)]

    subplot_kw = {
        'polar': True
    }

    fig, ax = plt.subplots(figsize=(8, 8), subplot_kw=subplot_kw)
    ax.set_theta_offset(np.pi / 2)
```

```

ax.set_theta_direction(-1)
ax.set_rlabel_position(0)

plt.xticks(angles[:-1], datas)
plt.yticks(color="grey", size=10)


# ADD PLOTS
for outcome in data_counts.columns:
    counts = data_counts[outcome].tolist()
    counts += counts[:1] # Properly loops the circle back
    ax.plot(angles, counts, linewidth=1, linestyle='solid', label=outcome)
    ax.fill(angles, counts, alpha=0.1)


plt.title(f"Counts by {column} Bins")
plt.legend(loc='upper right', bbox_to_anchor=(0.1, 0.1))

plt.show()


```

```
plot_radar(df, 20, "BMI", "Outcome")
```






```
plot_radar(df, 30, "Glucose", "Outcome")
```







```
# Outlier Detection:
outliers_cols = ["Glucose","BloodPressure","Insulin","BMI"]
outliers_df = pd.DataFrame()

for col in outliers_cols:
    stats = df[col].describe()

    IQR = stats['75%'] - stats['25%']
    upper_limit = stats['75%'] + 1.5 * IQR
    lower_limit = stats['25%'] - 1.5 * IQR

    outliers = df[(df[col] > upper_limit) | (df[col] < lower_limit)]

    if outliers.empty:
        print(f'\n No outlier found in: {col}')
    else:
        print(f'\n ----- {col} -----')
        outliers_df = pd.concat([outliers_df, outliers])
        print(f"Number of Outliers: {len(outliers)}")
        print(f'\n IQR: {IQR}')
        print(f'\n Outliers Lower Bound: {lower_limit}')
        print(f'\n Outliers Upper Bound: {upper_limit}')

print(f'Total number of outlier: {len(outliers_df)}')
```

```
No outlier found in: Glucose  
Number of Outliers: 0
```

```
IQR: 40.0
```

```
Outliers Lower Bound: 40.0
```

```
Outliers Upper Bound: 200.0
```

```
----- BloodPressure -----  
Number of Outliers: 10
```

```
IQR: 16.0
```

```
Outliers Lower Bound: 40.0
```

```
Outliers Upper Bound: 104.0
```

```
No outlier found in: Insulin  
Number of Outliers: 0
```

```
IQR: 67.0
```

```
Outliers Lower Bound: 2.0
```

```
Outliers Upper Bound: 270.0
```

```
----- BMI -----  
Number of Outliers: 4
```

```
IQR: 9.0
```

```
Outliers Lower Bound: 14.0
```

```
Outliers Upper Bound: 50.0  
Total number of outlier: 14
```

```
# Z-score determine:
all_outliers = pd.DataFrame()
# Compute mean and standard deviation
for col in outliers_cols:
    mean = df[col].mean()
    std_dev = df[col].std()

    # Calculating z-scores
    df['Z_Score'] = (df[col] - mean) / std_dev

    # outliers threshold
    threshold = 2

    # Identify outliers
    outliers = df[abs(df['Z_Score']) > threshold]
    all_outliers = pd.concat([all_outliers,outliers])

    print(len(outliers),col)
print(f'Total number of outlier: {len(all_outliers)}')
```

38 Glucose  
35 BloodPressure  
59 Insulin  
29 BMI  
Total number of outlier: 161

```
df.corr()['Outcome'] #other features correlation to Outcome
```

```
Pregnancies      0.296356
Glucose           0.494213
BloodPressure     0.182309
SkinThickness     0.281523
Insulin           0.509097
BMI               0.299717
DiabetesPedigreeFunction 0.166545
Age               0.246939
Outcome           1.000000
NewBMI_Obesity 1      NaN
NewBMI_Obesity 2     -0.155184
NewBMI_Obesity 3      0.155184
NewBMI_Overweight      NaN
NewBMI_Underweight     NaN
NewInsulinScore_Normal -0.650555
NewGlucose_Low         -0.088157
NewGlucose_Normal      -0.317394
NewGlucose_Overweight -0.119731
NewGlucose_Secret       0.419215
Z_Score               0.299717
Name: Outcome, dtype: float64
```

```
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(X)
X.shape
```

```
(760, 19)
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
```

```
import tensorflow
from tensorflow import keras
from tensorflow.keras import Sequential
from keras.layers import Dense,Dropout
```

```

model = Sequential()
model.add(Dense(32,activation='relu',input_dim=19))
model.add(Dense(1,activation='sigmoid'))

model.compile(optimizer='Adam',loss='binary_crossentropy',metrics=['accuracy'])

model.fit(X_train,y_train,batch_size=32,epochs=100,validation_data=(X_test,y_te

19/19 [=====] - 0s 17ms/step - loss: -178.7066 - a
Epoch 72/100
19/19 [=====] - 0s 10ms/step - loss: -183.7772 - a
Epoch 73/100
19/19 [=====] - 0s 17ms/step - loss: -188.6419 - a
Epoch 74/100
19/19 [=====] - 0s 7ms/step - loss: -193.8483 - ac
Epoch 75/100
19/19 [=====] - 0s 7ms/step - loss: -199.0373 - ac
Epoch 76/100
19/19 [=====] - 0s 6ms/step - loss: -204.3369 - ac
Epoch 77/100
19/19 [=====] - 0s 7ms/step - loss: -209.4307 - ac
Epoch 78/100
19/19 [=====] - 0s 6ms/step - loss: -214.9252 - ac
Epoch 79/100
19/19 [=====] - 0s 8ms/step - loss: -220.3127 - ac
Epoch 80/100
19/19 [=====] - 0s 7ms/step - loss: -225.8326 - ac
Epoch 81/100
19/19 [=====] - 0s 7ms/step - loss: -231.2605 - ac
Epoch 82/100
19/19 [=====] - 0s 5ms/step - loss: -236.7471 - ac
Epoch 83/100
19/19 [=====] - 0s 5ms/step - loss: -242.4043 - ac
Epoch 84/100
19/19 [=====] - 0s 6ms/step - loss: -248.0512 - ac
Epoch 85/100
19/19 [=====] - 0s 9ms/step - loss: -253.8454 - ac
Epoch 86/100
19/19 [=====] - 0s 6ms/step - loss: -259.4611 - ac
Epoch 87/100
19/19 [=====] - 0s 6ms/step - loss: -265.3672 - ac
Epoch 88/100
19/19 [=====] - 0s 7ms/step - loss: -271.2856 - ac
Epoch 89/100
19/19 [=====] - 0s 6ms/step - loss: -277.1169 - ac
Epoch 90/100
19/19 [=====] - 0s 8ms/step - loss: -283.1459 - ac
Epoch 91/100
19/19 [=====] - 0s 5ms/step - loss: -289.1490 - ac
Epoch 92/100
19/19 [=====] - 0s 8ms/step - loss: -295.0878 - ac

```

```

Epoch 93/100
19/19 [=====] - 0s 6ms/step - loss: -301.5058 - ac
Epoch 94/100
19/19 [=====] - 0s 7ms/step - loss: -307.5862 - ac
Epoch 95/100
19/19 [=====] - 0s 8ms/step - loss: -313.8767 - ac
Epoch 96/100
19/19 [=====] - 0s 6ms/step - loss: -320.1636 - ac
Epoch 97/100
19/19 [=====] - 0s 8ms/step - loss: -326.5222 - ac
Epoch 98/100
19/19 [=====] - 0s 7ms/step - loss: -332.7726 - ac
Epoch 99/100
19/19 [=====] - 0s 15ms/step - loss: -339.4695 - a
Epoch 100/100
19/19 [=====] - 0s 17ms/step - loss: -345.9737 - a
keras.callbacks.History at 0x7e013e5dd360

```

```
model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 32)	640
dense_5 (Dense)	(None, 1)	33
Total params: 673 (2.63 KB)		
Trainable params: 673 (2.63 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
model.fit(X_train,y_train,batch_size=32,epochs=100,initial_epoch=6,validation_c
```

```

Epoch 7/100
19/19 [=====] - 0s 12ms/step - loss: -359.5589 - a
Epoch 8/100
19/19 [=====] - 0s 8ms/step - loss: -366.1521 - ac
Epoch 9/100
19/19 [=====] - 0s 7ms/step - loss: -373.0780 - ac
Epoch 10/100
19/19 [=====] - 0s 6ms/step - loss: -379.9747 - ac
Epoch 11/100
19/19 [=====] - 0s 7ms/step - loss: -386.5775 - ac
Epoch 12/100
19/19 [=====] - 0s 9ms/step - loss: -393.6030 - ac
Epoch 13/100
19/19 [=====] - 0s 7ms/step - loss: -400.1753 - ac
Epoch 14/100
19/19 [=====] - 0s 7ms/step - loss: -407.5401 - ac
Epoch 15/100
19/19 [=====] - 0s 7ms/step - loss: -414.5459 - ac
Epoch 16/100
19/19 [=====] - 0s 6ms/step - loss: -421.6665 - ac
Epoch 17/100

```

```

Epoch 17/100
19/19 [=====] - 0s 8ms/step - loss: -428.6198 - ac
Epoch 18/100
19/19 [=====] - 0s 7ms/step - loss: -435.9584 - ac
Epoch 19/100
19/19 [=====] - 0s 9ms/step - loss: -443.1318 - ac
Epoch 20/100
19/19 [=====] - 0s 6ms/step - loss: -450.3473 - ac
Epoch 21/100
19/19 [=====] - 0s 7ms/step - loss: -457.7197 - ac
Epoch 22/100
19/19 [=====] - 0s 5ms/step - loss: -465.2077 - ac
Epoch 23/100
19/19 [=====] - 0s 6ms/step - loss: -472.4452 - ac
Epoch 24/100
19/19 [=====] - 0s 5ms/step - loss: -479.8652 - ac
Epoch 25/100
19/19 [=====] - 0s 6ms/step - loss: -487.6178 - ac
Epoch 26/100
19/19 [=====] - 0s 6ms/step - loss: -495.0153 - ac
Epoch 27/100
19/19 [=====] - 0s 6ms/step - loss: -502.6337 - ac
Epoch 28/100
19/19 [=====] - 0s 6ms/step - loss: -510.2431 - ac
Epoch 29/100
19/19 [=====] - 0s 6ms/step - loss: -517.8948 - ac
Epoch 30/100
19/19 [=====] - 0s 6ms/step - loss: -525.6293 - ac
Epoch 31/100
19/19 [=====] - 0s 6ms/step - loss: -533.4838 - ac
Epoch 32/100
19/19 [=====] - 0s 5ms/step - loss: -541.3406 - ac
Epoch 33/100
19/19 [=====] - 0s 4ms/step - loss: -549.1597 - ac
Epoch 34/100
19/19 [=====] - 0s 4ms/step - loss: -556.9739 - ac
Epoch 35/100
19/19 [=====] - 0s 4ms/step - loss: -565.1715 - ac
Epoch 36/100
19/19 [=====] - 0s 4ms/step - loss: -573.0000 - ac

```

```
def build_model(hp):
```

```
    model = Sequential()
```

```
    units = hp.Int('units',min_value = 8,max_value = 128,step = 8)
```

```
model.add(Dense(units=units,activation='relu',input_dim=19))
model.add(Dense(1,activation='sigmoid'))
```

```
model.compile(optimizer='rmsprop',loss='binary_crossentropy', metrics=['acc'])
return model
```

```
model.fit(X_train,y_train,batch_size=32, epochs=100, initial_epoch=5,validation_data=(X_val,y_val))
```

```
Epoch 71/100
19/19 [=====] - 0s 4ms/step - loss: -1965.0654 - a
Epoch 72/100
19/19 [=====] - 0s 5ms/step - loss: -1977.7638 - a
Epoch 73/100
19/19 [=====] - 0s 5ms/step - loss: -1991.3127 - a
Epoch 74/100
19/19 [=====] - 0s 4ms/step - loss: -2004.8585 - a
Epoch 75/100
19/19 [=====] - 0s 3ms/step - loss: -2017.7714 - a
Epoch 76/100
19/19 [=====] - 0s 4ms/step - loss: -2031.3949 - a
Epoch 77/100
19/19 [=====] - 0s 4ms/step - loss: -2044.4111 - a
Epoch 78/100
19/19 [=====] - 0s 5ms/step - loss: -2058.3169 - a
Epoch 79/100
19/19 [=====] - 0s 3ms/step - loss: -2071.3398 - a
Epoch 80/100
19/19 [=====] - 0s 4ms/step - loss: -2085.1768 - a
Epoch 81/100
19/19 [=====] - 0s 4ms/step - loss: -2098.5354 - a
Epoch 82/100
19/19 [=====] - 0s 3ms/step - loss: -2112.4033 - a
Epoch 83/100
19/19 [=====] - 0s 3ms/step - loss: -2125.6433 - a
Epoch 84/100
19/19 [=====] - 0s 4ms/step - loss: -2139.7600 - a
Epoch 85/100
19/19 [=====] - 0s 5ms/step - loss: -2152.7939 - a
Epoch 86/100
19/19 [=====] - 0s 6ms/step - loss: -2166.9785 - a
Epoch 87/100
19/19 [=====] - 0s 7ms/step - loss: -2180.7422 - a
Epoch 88/100
19/19 [=====] - 0s 7ms/step - loss: -2194.5342 - a
Epoch 89/100
19/19 [=====] - 0s 5ms/step - loss: -2208.1719 - a
Epoch 90/100
19/19 [=====] - 0s 6ms/step - loss: -2222.2021 - a
Epoch 91/100
19/19 [=====] - 0s 5ms/step - loss: -2236.0034 - a
Epoch 92/100
19/19 [=====] - 0s 7ms/step - loss: -2249.9890 - a
Epoch 93/100
19/19 [=====] - 0s 5ms/step - loss: -2264.0325 - a
Epoch 94/100
19/19 [=====] - 0s 6ms/step - loss: -2277.7388 - a
Epoch 95/100
```



```
Epoch 95/100
19/19 [=====] - 0s 5ms/step - loss: -2291.8398 - a
Epoch 96/100
19/19 [=====] - 0s 6ms/step - loss: -2306.1133 - a
Epoch 97/100
19/19 [=====] - 0s 6ms/step - loss: -2319.8721 - a
Epoch 98/100
19/19 [=====] - 0s 7ms/step - loss: -2334.0767 - a
Epoch 99/100
19/19 [=====] - 0s 5ms/step - loss: -2348.3430 - a
Epoch 100/100
19/19 [=====] - 0s 5ms/step - loss: -2362.2266 - a
```

```
model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 32)	640
dense_5 (Dense)	(None, 1)	33
Total params: 673 (2.63 KB)		
Trainable params: 673 (2.63 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
def build_model(hp):
```

```
    model = Sequential()
```

```
    model.add(Dense(64,activation='relu',input_dim=19))
```

```
    for i in range(hp.Int('num_layers',min_value=1,max_value=10)):
```

```
        model.add(Dense(64,activation='relu'))
```

```
    model.add(Dense(1,activation='sigmoid'))
```

```
    model.compile(optimizer='rmsprop', loss='binary_crossentropy',metrics=['acc
```

```
    return model
```

```
model.fit(X_train,y_train,epochs=100,initial_epoch=6,validation_data=(X_test,y_
```

```
19/19 [=====] - 0s 12ms/step - loss: -5061.8696 -
Epoch 72/100
19/19 [=====] - 0s 12ms/step - loss: -5081.4297 -
Epoch 73/100
19/19 [=====] - 0s 11ms/step - loss: -5101.0430 -
Epoch 74/100
19/19 [=====] - 0s 9ms/step - loss: -5121.1108 - a
Epoch 75/100
19/19 [=====] - 0s 13ms/step - loss: -5141.2559 -
```

```

Epoch 76/100
19/19 [=====] - 0s 10ms/step - loss: -5160.4673 - a
Epoch 77/100
19/19 [=====] - 0s 7ms/step - loss: -5180.8267 - a
Epoch 78/100
19/19 [=====] - 0s 8ms/step - loss: -5200.9785 - a
Epoch 79/100
19/19 [=====] - 0s 9ms/step - loss: -5220.7808 - a
Epoch 80/100
19/19 [=====] - 0s 8ms/step - loss: -5241.3086 - a
Epoch 81/100
19/19 [=====] - 0s 11ms/step - loss: -5261.0015 - a
Epoch 82/100
19/19 [=====] - 0s 9ms/step - loss: -5281.1011 - a
Epoch 83/100
19/19 [=====] - 0s 7ms/step - loss: -5301.3477 - a
Epoch 84/100
19/19 [=====] - 0s 11ms/step - loss: -5321.8037 - a
Epoch 85/100
19/19 [=====] - 0s 7ms/step - loss: -5341.5508 - a
Epoch 86/100
19/19 [=====] - 0s 10ms/step - loss: -5361.9570 - a
Epoch 87/100
19/19 [=====] - 0s 11ms/step - loss: -5382.6289 - a
Epoch 88/100
19/19 [=====] - 0s 11ms/step - loss: -5402.7012 - a
Epoch 89/100
19/19 [=====] - 0s 16ms/step - loss: -5423.1416 - a
Epoch 90/100
19/19 [=====] - 0s 11ms/step - loss: -5443.2788 - a
Epoch 91/100
19/19 [=====] - 0s 7ms/step - loss: -5463.9888 - a
Epoch 92/100
19/19 [=====] - 0s 7ms/step - loss: -5484.6450 - a
Epoch 93/100
19/19 [=====] - 0s 8ms/step - loss: -5504.8940 - a
Epoch 94/100
19/19 [=====] - 0s 7ms/step - loss: -5525.3643 - a
Epoch 95/100
19/19 [=====] - 0s 8ms/step - loss: -5545.8008 - a
Epoch 96/100
19/19 [=====] - 0s 8ms/step - loss: -5566.6001 - a
Epoch 97/100
19/19 [=====] - 0s 13ms/step - loss: -5587.5200 - a
Epoch 98/100
19/19 [=====] - 0s 11ms/step - loss: -5607.8042 - a
Epoch 99/100
19/19 [=====] - 0s 7ms/step - loss: -5628.7319 - a
Epoch 100/100
19/19 [=====] - 0s 7ms/step - loss: -5649.7192 - a
<keras.src.callbacks.History at 0x7e013a399840>

```

```
model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```
=====
dense_4 (Dense)                (None, 32)                640
dense_5 (Dense)                (None, 1)                 33
=====
Total params: 673 (2.63 KB)
Trainable params: 673 (2.63 KB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
def build_model(hp):

    model = Sequential()
    counter = 0

    for i in range(hp.Int('num_layers',min_value = 1,max_value = 10)):
        if counter == 0:
            model.add(
                Dense(
                    hp.Int('units' + str(i),min_value = 8,max_value = 128,step
                    activation= hp.Choice('activation' + str(i), values=['relu'
                    input_dim=8
                    )
                )
            model.add(Dropout(hp.Choice('dropout'+ str(i), values=[0.1,0.2,0.3,
        else:
            model.add(
                Dense(
                    hp.Int('units' + str(i),min_value = 8,max_value = 128,step
                    activation= hp.Choice('activation' + str(i), values=['relu'
                    )
                )
            model.add(Dropout(hp.Choice('dropout'+ str(i), values=[0.1,0.2,0.3,
        counter += 1

    model.add(Dense(1,activation='sigmoid'))

    model.compile(optimizer=hp.Choice('optimizer', values=['rmsprop','adam','sg
                    loss='binary_crossentropy',
                    metrics=['accuracy']))

    return model
model.fit(X_train,y_train,epochs=200, initial_epoch=5,validation_data=(X_test,y

Epoch 171/200
19/19 [=====] - 0s 6ms/step - loss: -9557.0674 - a
Epoch 172/200
19/19 [=====] - 0s 16ms/step - loss: -9583.4990 -
Epoch 173/200
19/19 [=====] - 0s 8ms/step - loss: -9609.4092 - a
Epoch 174/200
19/19 [=====] - 0s 6ms/step - loss: -9636.5801 - a
Epoch 175/200
```

```

19/19 [=====] - 0s 8ms/step - loss: -9662.4141 - a
Epoch 176/200
19/19 [=====] - 0s 11ms/step - loss: -9689.2002 -
Epoch 177/200
19/19 [=====] - 0s 6ms/step - loss: -9716.1211 - a
Epoch 178/200
19/19 [=====] - 0s 6ms/step - loss: -9742.0820 - a
Epoch 179/200
19/19 [=====] - 0s 8ms/step - loss: -9768.8848 - a
Epoch 180/200
19/19 [=====] - 0s 8ms/step - loss: -9795.3506 - a
Epoch 181/200
19/19 [=====] - 0s 11ms/step - loss: -9822.0537 -
Epoch 182/200
19/19 [=====] - 0s 7ms/step - loss: -9849.6113 - a
Epoch 183/200
19/19 [=====] - 0s 8ms/step - loss: -9875.5967 - a
Epoch 184/200
19/19 [=====] - 0s 6ms/step - loss: -9902.2588 - a
Epoch 185/200
19/19 [=====] - 0s 6ms/step - loss: -9929.0010 - a
Epoch 186/200
19/19 [=====] - 0s 6ms/step - loss: -9956.2715 - a
Epoch 187/200
19/19 [=====] - 0s 7ms/step - loss: -9982.9941 - a
Epoch 188/200
19/19 [=====] - 0s 8ms/step - loss: -10009.8203 -
Epoch 189/200
19/19 [=====] - 0s 7ms/step - loss: -10037.3955 -
Epoch 190/200
19/19 [=====] - 0s 7ms/step - loss: -10064.0557 -
Epoch 191/200
19/19 [=====] - 0s 9ms/step - loss: -10091.4014 -
Epoch 192/200
19/19 [=====] - 0s 7ms/step - loss: -10117.2627 -
Epoch 193/200
19/19 [=====] - 0s 22ms/step - loss: -10145.4756 -
Epoch 194/200
19/19 [=====] - 0s 13ms/step - loss: -10172.5273 -
Epoch 195/200
19/19 [=====] - 0s 12ms/step - loss: -10200.1123 -
Epoch 196/200
19/19 [=====] - 0s 12ms/step - loss: -10226.5576 -
Epoch 197/200
19/19 [=====] - 0s 12ms/step - loss: -10253.8857 -
Epoch 198/200
19/19 [=====] - 0s 24ms/step - loss: -10281.0830 -
Epoch 199/200
19/19 [=====] - 1s 33ms/step - loss: -10308.6787 -
Epoch 200/200
19/19 [=====] - 1s 40ms/step - loss: -10336.1689 -

```

```
model.summary()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
=====		

dense_4 (Dense)	(None, 32)	640
dense_5 (Dense)	(None, 1)	33

```
=====
Total params: 673 (2.63 KB)
Trainable params: 673 (2.63 KB)
Non-trainable params: 0 (0.00 Byte)
```

---